

Dev1num3

February 5, 2019

```
In [1]: import torch
import torchvision
import torchvision.transforms as transforms #
import torchvision.datasets.mnist as mnist # to import data

# we use torch.cuda.Event(enable_timing=True) to measure time
# from timeit import default_timer as timer
# import time

import torch.optim as optim
import torch.nn as nn
import torch.nn.functional as F

import collections          # for ordered_dictionnary
import torch.nn.init as init # to initialize model

import copy                  # for copy.deepcopy( ... )

import matplotlib.pyplot as plt
import numpy as np

from __future__ import print_function, division
import os
from PIL import Image
import torch
import pandas as pd
from skimage import io, transform
import numpy as np
import matplotlib.pyplot as plt
from torch.utils.data import Dataset, DataLoader
from torchvision import transforms, utils

import collections
import torch.nn.init as init

In [5]: total_nb_of_sample = 9999 # total number of sample per class, this shouldn't be hardcoded
class labeledDataSet(torch.utils.data.Dataset):
```

```

def __init__(self, label, idx_min , idx_max , root_dir ):
    """
    Args:
        label is either "Cat" or "Dog"
        load in the dataset picture no. idx_min to idx_max included
        root_dir(string): directory with all images with the same label
    """
    super(labeledDataSet, self).__init__()
    dic = {"Cat" : int(1) , "Dog" : int(0)}

    if dic.get(label , "not found" ) == "not found" :
        raise ValueError("label must either be Cat or Dog")

    if idx_max > total_nb_of_sample or idx_min < 0 :
        raise ValueError("min, max index error")

    self.label_name = label
    self.idx_min = idx_min
    self.idx_max = idx_max
    self.label = dic[label]
    self.root_dir = root_dir
    self.load_data()

def load_data(self) :
    size = self.__len__()

    self.data_tensor = torch.empty(size,3,64,64,dtype=torch.float)
    self.target_tensor = torch.ones(size,dtype=torch.long) * self.label

    for i,j in enumerate(range( self.idx_min , self.idx_max , 1 ),0) :
        img_name = os.path.join(self.root_dir, "{index}.{Label}.jpg".format(index=
        # image = io.imread(img_name)
        # print( i , " : " , io.imread(img_name).shape )
        # img = io.imread(img_name)
        img = Image.open( img_name ).convert('RGB')
        """
        NOTE : There is at least one grey-scale picture of Einstein, manually remove
        the greyscale (1 channel) images have to be converted to 3 channels tensor.
        """
        image = torch.from_numpy( np.transpose( img , (-1,-3,-2) ) )
        self.data_tensor[i,:,:,:] = image

def __len__(self):
    # return len(self.name_frame)
    return (self.idx_max - self.idx_min + 1)

def __getitem__(self, idx):
    return self.data_tensor[idx], self.target_tensor[idx]

```

Setting the directory of the pictures

```
In [ ]: dog_dir = "./data_catdogs/trainset/Dog"
        cat_dir = "./data_catdogs/trainset/Cat"

        # picture no 1          to (idx_sep)          belong to the training dataset
        # picture no (idx_sep + 1) to (total_nb_of_sample) belong to the validation dataset
        idx_sep = 8000

        trainDogSet = labeledDataSet( "Dog", 1 , idx_sep , root_dir = dog_dir )
        trainCatSet = labeledDataSet( "Cat", 1 , idx_sep , root_dir = cat_dir )

        testsDogSet = labeledDataSet( "Dog", idx_sep + 1 , total_nb_of_sample , root_dir = dog_dir )
        testsCatSet = labeledDataSet( "Cat", idx_sep + 1 , total_nb_of_sample , root_dir = cat_dir )

In [6]: trainset = torch.utils.data.ConcatDataset( [ trainDogSet , trainCatSet ] )
        testsset = torch.utils.data.ConcatDataset( [ testsDogSet , testsCatSet ] )

In [7]: testing_dataset_size = testsCatSet.__len__() + testsDogSet.__len__()
        training_dataset_size = trainCatSet.__len__() + trainDogSet.__len__()

        print( "training dataset size = " , training_dataset_size )
        print( "testing dataset size = " , testing_dataset_size )
        print( trainDogSet.__doc__ )
        img, label = trainDogSet.__getitem__(1)
        print( "img size = " , img.size() , "label size = " , label.size() )

training dataset size = 16000
testing dataset size = 3998
None
img = torch.Size([3, 64, 64]) label = torch.Size([])
```

Display some samples

```
In [8]: nb_sample = 8

        trainloader = torch.utils.data.DataLoader(trainset , batch_size = nb_sample, shuffle=True)
        testsloader = torch.utils.data.DataLoader(testsset , batch_size = nb_sample, shuffle=True)

        # functions to show an image

        def imshow(img):
            npimg = img.numpy() / 255
            plt.imshow(np.transpose(npimg, (1, 2, 0)))
            plt.show()

        # get some random training images
        dataiter = iter( testsloader )
```

```

images, labels = dataiter.next()

# show images
imshow(torchvision.utils.make_grid(images))
# print labels
print(' '.join('%5s' % labels[j].item() for j in range(nb_sample)))

```



```

0    0    1    1    0    0    0    0

```

set the device

```

In [9]: device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
        print(device)

```

cuda:0

0.1 The model

architecture taken from : https://github.com/MaximumEntropy/welcome_tutorials/tree/pytorch/pytorch

```

In [10]: class Classifier(nn.Module):
        """Convnet Classifier"""
        def __init__(self, kernel_sz = 3 ):

            if kernel_sz % 2 == 0 :
                raise ValueError("kernel size must be odd")
            pad = kernel_sz // 2

            super(Classifier, self).__init__()
            self.conv = nn.Sequential(
                # Layer 1
                nn.Conv2d(in_channels=3, out_channels=16, kernel_size=( kernel_sz , kernel_sz ),
                           padding=pad),
                # nn.Dropout(p=0.5),
                nn.ReLU(),
                nn.MaxPool2d(kernel_size=(2, 2), stride=2),

                # Layer 2

```

```

nn.Conv2d(in_channels=16, out_channels=32, kernel_size=( kernel_sz , kernel_sz ),
# nn.Dropout(p=0.5),
nn.ReLU(),
nn.MaxPool2d(kernel_size=(2, 2), stride=2),

# Layer 3
nn.Conv2d(in_channels=32, out_channels=64, kernel_size=( kernel_sz , kernel_sz ),
# nn.Dropout(p=0.5),
nn.ReLU(),
nn.MaxPool2d(kernel_size=(2, 2), stride=2),

# Layer 4
nn.Conv2d(in_channels=64, out_channels=128, kernel_size=(kernel_sz , kernel_sz ),
# nn.Dropout(p=0.5),
nn.ReLU(),
nn.MaxPool2d(kernel_size=(2, 2), stride=2)
)
#
self.fct = nn.Linear(4*4*128, 1)

def forward(self, x):
x = self.conv(x)
x = x.view(-1, 4*4*128)
x = torch.sigmoid(self.fct(x))
# x = self.fct(x)
# x = F.softmax(self.fct(x), dim=-1)
return x

```

Print the number of parameters in each models and display the computation

```

In [11]: def number_of_params( net ) :
nb_param = 0
param_lst = " "
for i, (key, value) in enumerate( net.state_dict().items() ) :

    if i == 0 :
        param_lst = param_lst + "\n ({:<20}      ".format( key + " )" )
    else :
        param_lst = param_lst + "\n ({:<20}  + ".format( key + " )" )

nb_param_tmp = 1

for j , x in enumerate(value.size()) :
    if j == 0 :
        param_lst = param_lst + "{xx}".format( xx = x )
    else :
        param_lst = param_lst + "*{xx}".format( xx = x )

```

```

        nb_param_tmp = nb_param_tmp * x

    nb_param = nb_param + nb_param_tmp

    print( "number of params = " , nb_param , " = ", param_lst )

mytestnet = Classifier( kernel_sz=9 )
number_of_params( mytestnet )
del mytestnet

number of params = 877089 =
(conv.0.weight)      16*3*9*9
(conv.0.bias)         + 16
(conv.3.weight)       + 32*16*9*9
(conv.3.bias)         + 32
(conv.6.weight)       + 64*32*9*9
(conv.6.bias)         + 64
(conv.9.weight)       + 128*64*9*9
(conv.9.bias)         + 128
(fct.weight)          + 1*2048
(fct.bias)            + 1

init method

In [12]: def gloriot_init ( layer ) :
        """
        Weights are generated from  $U[-d,d]$  where  $d = \sqrt{6/(fan\_in + fan\_out)}$ , biases are
        """
        if type(layer) == nn.Linear or type(layer) == nn.Conv2d :
            init.xavier_uniform_( layer.weight , gain=1 )
            layer.bias.data.fill_(0.0)

In [33]: if True :
        del cudanet

In [34]: cudanet = Classifier( kernel_sz = 9 )
        cudanet.apply( gloriot_init )
        cudanet.to(device)

Out[34]: Classifier(
  (conv): Sequential(
    (0): Conv2d(3, 16, kernel_size=(9, 9), stride=(1, 1), padding=(4, 4))
    (1): ReLU()
    (2): MaxPool2d(kernel_size=(2, 2), stride=2, padding=0, dilation=1, ceil_mode=False)
    (3): Conv2d(16, 32, kernel_size=(9, 9), stride=(1, 1), padding=(4, 4))
    (4): ReLU()
    (5): MaxPool2d(kernel_size=(2, 2), stride=2, padding=0, dilation=1, ceil_mode=False)

```

```

(6): Conv2d(32, 64, kernel_size=(9, 9), stride=(1, 1), padding=(4, 4))
(7): ReLU()
(8): MaxPool2d(kernel_size=(2, 2), stride=2, padding=0, dilation=1, ceil_mode=False)
(9): Conv2d(64, 128, kernel_size=(9, 9), stride=(1, 1), padding=(4, 4))
(10): ReLU()
(11): MaxPool2d(kernel_size=(2, 2), stride=2, padding=0, dilation=1, ceil_mode=False)
)
(fct): Linear(in_features=2048, out_features=1, bias=True)
)

```

```

In [36]: criterion = nn.BCELoss() # Binary Cross Entropy
optimizer = optim.SGD(cudanet.parameters(), lr=0.0001, momentum=0.0, weight_decay=0)
nb_epoch = 30

```

```

In [37]: if False :
    print( cudanet( torch.ones(1,3,64,64).to(device) ) )
    print( cudanet( torch.ones(1,3,64,64).to(device) ).dtype )

    with torch.no_grad() :
        for i, data in enumerate(trainloader, 0):
            # get the inputs
            inputs, labels = data
            labels = labels.float()
            inputs, labels = inputs.to(device), labels.to(device)

            # zero the parameter gradients
            optimizer.zero_grad()

            # forward + backward + optimize
            outputs = cudanet(inputs).squeeze()
            print( outputs.size() , labels.size() )
            print( outputs )
            loss = criterion(outputs, labels)
            print( loss )

            break

```

0.2 Training the model

Note : 30 epoch with lr=0.0001 are not enough, add more epoch, larger lr could be an issue for fine tuning the model

```

In [39]: trainloader = torch.utils.data.DataLoader(trainset, batch_size=4*64,shuffle=True, num_workers=4,
state_dict_list = list())

torch.cuda.synchronize()

start = torch.cuda.Event(enable_timing=True)
end = torch.cuda.Event(enable_timing=True)

```

```

start.record()

for epoch in range( nb_epoch ):  # loop over the dataset multiple times

    running_loss = 0.0
    for i, data in enumerate(trainloader, 0):
        # get the inputs
        inputs, labels = data
        labels = labels.float()
        inputs, labels = inputs.to(device), labels.to(device)

        # zero the parameter gradients
        optimizer.zero_grad()

        # forward + backward + optimize
        outputs = cudanet(inputs).squeeze()
        # if i == 0 : print( outputs )
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        # print statistics
        running_loss += loss.item()
    else : # print every epoch
        print('epoch = %d, loss = %.8f' % (epoch + 1, running_loss / (i*8*64))) # nb
        running_loss = 0.0
        torch.cuda.synchronize()
        tmp_state_dict = {}
        for k, v in cudanet.state_dict().items():
            tmp_state_dict[k] = v.cpu()
        state_dict_list.append( tmp_state_dict )
        torch.cuda.synchronize()

else :
    print('Finished Training')

end.record()
torch.cuda.synchronize()
print( "time required = " , start.elapsed_time(end)*0.001 , " s ")

```

```

epoch = 1, loss = 0.00791041
epoch = 2, loss = 0.00150025
epoch = 3, loss = 0.00140869
epoch = 4, loss = 0.00135263
epoch = 5, loss = 0.00131762
epoch = 6, loss = 0.00128026
epoch = 7, loss = 0.00126457
epoch = 8, loss = 0.00125376

```



```

epoch = 9, loss = 0.00122598
epoch = 10, loss = 0.00122030
epoch = 11, loss = 0.00119520
epoch = 12, loss = 0.00118546
epoch = 13, loss = 0.00117974
epoch = 14, loss = 0.00116469
epoch = 15, loss = 0.00115254
epoch = 16, loss = 0.00115322
epoch = 17, loss = 0.00114271
epoch = 18, loss = 0.00113528
epoch = 19, loss = 0.00112319
epoch = 20, loss = 0.00111661
epoch = 21, loss = 0.00110833
epoch = 22, loss = 0.00110505
epoch = 23, loss = 0.00109424
epoch = 24, loss = 0.00109197
epoch = 25, loss = 0.00107874
epoch = 26, loss = 0.00108290
epoch = 27, loss = 0.00107191
epoch = 28, loss = 0.00106411
epoch = 29, loss = 0.00106355
epoch = 30, loss = 0.00104832
Finished Training
time required = 366.51428125 s

```

test accuracy of the current state on validation dataset

```

In [29]: testloader = torch.utils.data.DataLoader(testsset, batch_size=8*64,shuffle=False, num_workers=4)
correct = torch.tensor([0])
total = torch.tensor([0])

correct, total = correct.to(device) , total.to(device)
with torch.no_grad():
    for data in testloader:
        images, labels = data
        images, labels = images.to(device), labels.to(device)
        outputs = cudanet(images).squeeze()
        predicted = torch.where(
            outputs > 0.5, torch.ones_like(labels, device=device) , torch.zeros_like(labels, device=device)
        )
        total += labels.size(0)
        correct += (predicted == labels).sum()

print('Accuracy of the network on the' , testing_dataset_size , 'test images: %.2f %%' % ( (100 * correct.double()) / total.double() ) )

```

Accuracy of the network on the 3998 test images: 70.19 %

accuracy across epoch

```
In [40]: trainloader = torch.utils.data.DataLoader(trainset , batch_size = 8*64 , shuffle=False)
        testsloader = torch.utils.data.DataLoader(testsset , batch_size = 8*64, shuffle=False)

accuracy = torch.ones(nb_epoch,2, dtype=torch.float) * 100
for epoch , tmp_state_dict in enumerate(state_dict_list,0) :
    if epoch % 1 != 0 :
        continue

    cudaTOCPUNet = copy.deepcopy( cudanet )
    cudaTOCPUNet.load_state_dict( tmp_state_dict )
    cuda_test_net = copy.deepcopy(cudaTOCPUNet).to(device)

    correct = torch.tensor([0,0])
    total   = torch.tensor([0,0])

    correct, total = correct.to(device) , total.to(device)

    loader_list = [ testsloader , trainloader ]
    with torch.no_grad():
        for i, loader in enumerate(loader_list,0) :
            for data in loader:
                images, labels = data
                images, labels = images.to(device), labels.to(device)
                outputs = cuda_test_net(images).squeeze()
                predicted = torch.where(
                    outputs > 0.5, torch.ones_like(labels, device=device)
                )
                total[i] += labels.size(0)
                correct[i] += (predicted == labels).sum()

    accuracy[epoch,:] = accuracy[epoch,:] * correct.type(torch.FloatTensor) / total.type(torch.FloatTensor)

    print('epoch %3d : Accuracy of the network on the test images: %.2f %% , training images %.2f %%' % ( epoch+1, accuracy[epoch,0] , accuracy[epoch,1] )
          )

epoch 1 : Accuracy of the network on the test images: 54.23 % , training images 55.58 %
epoch 2 : Accuracy of the network on the test images: 55.85 % , training images 56.51 %
epoch 3 : Accuracy of the network on the test images: 61.83 % , training images 61.86 %
epoch 4 : Accuracy of the network on the test images: 61.56 % , training images 61.91 %
epoch 5 : Accuracy of the network on the test images: 64.66 % , training images 64.74 %
epoch 6 : Accuracy of the network on the test images: 63.78 % , training images 64.81 %
epoch 7 : Accuracy of the network on the test images: 64.88 % , training images 65.60 %
```

```

epoch 8 : Accuracy of the network on the test images: 65.48 % , training images 66.20 %
epoch 9 : Accuracy of the network on the test images: 65.71 % , training images 67.39 %
epoch 10 : Accuracy of the network on the test images: 65.96 % , training images 67.42 %
epoch 11 : Accuracy of the network on the test images: 66.33 % , training images 68.41 %
epoch 12 : Accuracy of the network on the test images: 67.06 % , training images 68.57 %
epoch 13 : Accuracy of the network on the test images: 66.78 % , training images 69.58 %
epoch 14 : Accuracy of the network on the test images: 67.13 % , training images 69.89 %
epoch 15 : Accuracy of the network on the test images: 67.21 % , training images 70.06 %
epoch 16 : Accuracy of the network on the test images: 67.76 % , training images 70.53 %
epoch 17 : Accuracy of the network on the test images: 67.11 % , training images 69.84 %
epoch 18 : Accuracy of the network on the test images: 66.88 % , training images 69.39 %
epoch 19 : Accuracy of the network on the test images: 68.28 % , training images 71.71 %
epoch 20 : Accuracy of the network on the test images: 68.08 % , training images 71.53 %
epoch 21 : Accuracy of the network on the test images: 68.48 % , training images 72.18 %
epoch 22 : Accuracy of the network on the test images: 68.66 % , training images 72.34 %
epoch 23 : Accuracy of the network on the test images: 62.83 % , training images 64.88 %
epoch 24 : Accuracy of the network on the test images: 67.06 % , training images 70.85 %
epoch 25 : Accuracy of the network on the test images: 68.63 % , training images 72.71 %
epoch 26 : Accuracy of the network on the test images: 68.83 % , training images 73.53 %
epoch 27 : Accuracy of the network on the test images: 65.56 % , training images 68.79 %
epoch 28 : Accuracy of the network on the test images: 68.61 % , training images 73.47 %
epoch 29 : Accuracy of the network on the test images: 67.68 % , training images 72.28 %
epoch 30 : Accuracy of the network on the test images: 69.63 % , training images 74.24 %

```

0.2.1 Plot accuracy

```

In [43]: from matplotlib.pyplot import figure
         # import warnings
         # warnings.filterwarnings('ignore')

         # data to plot
         n_groups = nb_epoch
         accuracy_toplot = accuracy.numpy()

         tests_accuracy = accuracy_toplot[:,0]
         train_accuracy = accuracy_toplot[:,1]

         #
         plt.rcParams.update({'font.size': 16})
         plt.rcParams["figure.figsize"] = (12 ,8)
         # create plot
         fig, ax = plt.subplots()
         index = np.arange(n_groups)
         bar_width = 0.3
         opacity = 0.8

         rects1 = plt.bar(index, tests_accuracy, bar_width,

```

```

        alpha=opacity,
        color='blue',
        label='Test set')

rects2 = plt.bar(index + bar_width, train_accuracy, bar_width,
                 alpha=opacity,
                 color='r',
                 label='Training set')


eps = 3
top = min(int( np.ceil(accuracy_toplot.max() + eps)) , 100)
bot = max(int(np.floor(accuracy_toplot.min() - eps)) , 0 )

plt.ylim(bot, top)      # set the ylim to bottom, top
plt.axhline(y=75,color="black")

plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.title('Comparison between training set and test set accuracy \nduring the training')
plt.xticks(index + bar_width, range(1,n_groups,1) )
plt.yticks( range(bot,top+1,1) )
plt.legend()

# plt.tight_layout()
plt.show()

```



1 Misc

print some test sample that the net misclassifies

```
In [60]: testloader = torch.utils.data.DataLoader(testsset , batch_size = 8*64, shuffle=True

nb_of_error = torch.tensor([8])
j = torch.tensor([0])
errorimages = torch.empty(8,3,64,64)
errorlabels = torch.empty(8)
erroroutputs = torch.empty(8)
for images, labels in testloader :
    images, labels = images.to(device), labels.to(device)
    outputs = cuda_test_net(images).squeeze()
    predicted = torch.where(
        outputs > 0.5, torch.ones_like(labels, device=device) , torch.zeros_like(labels, device=device)
    )
    if not (predicted == labels).sum() == labels.size()[0] :
        comparison = (predicted == labels)
        for i,b in enumerate( comparison ) :
            if b == 0 :
```


```

        errorimages[j,:,:,:] = copy.deepcopy(images[i,:,:,:]).cpu()
        errorlabels[j] = labels[i].clone().detach().requires_grad_(False).cpu()
        erroroutputs[j] = outputs[i].clone().detach().requires_grad_(False).cpu()
        j = j + 1
        if j.item() >= nb_of_error.item() :
            break
    else :
        continue
    break

if j.item() == 0 :
    print( "no error found")
else :
    imshow( torchvision.utils.make_grid(errorimages) )

    print( "Probabilities goes from : 0 -> this is a dog for sure , to 1 -> cat for sure" )
    print( "this should be          : " , ",".join( " %-4d" % nb.item() for nb in errorlabels ) )
    print( "net associated prob      : " , ",".join( "%5.2f" % nb.item() for nb in erroroutputs ) )

```



```

Probabilities goes from : 0 -> this is a dog for sure , to 1 -> cat for sure
this should be          :  0   , 0   , 1   , 1   , 1   , 1   , 0   , 0
net associated prob      :  0.60, 0.85, 0.47, 0.25, 0.14, 0.41, 0.54, 0.76

```

1.0.1 Save and load models

Save the state_dict of the model for each epoch on a local directory

```

In [57]: local_path = "./save"
         for epoch , tmp_state_dict in enumerate(state_dict_list,1) :
             saving_path = local_path + "/dev1num3model_for_epoch{Epoch}.pth".format( Epoch = epoch )
             torch.save( tmp_state_dict , saving_path )
             # print( saving_path )

./save/dev1num3model_for_epoch1.pth
./save/dev1num3model_for_epoch2.pth
./save/dev1num3model_for_epoch3.pth
./save/dev1num3model_for_epoch4.pth
./save/dev1num3model_for_epoch5.pth

```

```

./save/dev1num3model_for_epoch6.pth
./save/dev1num3model_for_epoch7.pth
./save/dev1num3model_for_epoch8.pth
./save/dev1num3model_for_epoch9.pth
./save/dev1num3model_for_epoch10.pth
./save/dev1num3model_for_epoch11.pth
./save/dev1num3model_for_epoch12.pth
./save/dev1num3model_for_epoch13.pth
./save/dev1num3model_for_epoch14.pth
./save/dev1num3model_for_epoch15.pth
./save/dev1num3model_for_epoch16.pth
./save/dev1num3model_for_epoch17.pth
./save/dev1num3model_for_epoch18.pth
./save/dev1num3model_for_epoch19.pth
./save/dev1num3model_for_epoch20.pth
./save/dev1num3model_for_epoch21.pth
./save/dev1num3model_for_epoch22.pth
./save/dev1num3model_for_epoch23.pth
./save/dev1num3model_for_epoch24.pth
./save/dev1num3model_for_epoch25.pth
./save/dev1num3model_for_epoch26.pth
./save/dev1num3model_for_epoch27.pth
./save/dev1num3model_for_epoch28.pth
./save/dev1num3model_for_epoch29.pth
./save/dev1num3model_for_epoch30.pth

```

load from file and set the load the state_dict of the last epoch on a object the files have to be located in `./save` and named `"dev1num3model_for_epoch{j}.pth"` for `j` from ... to ...

```

In [59]: local_path = "./save"
        state_dict_list = list()

        from_idx = 1
        to_idx    = nb_epoch

        for epoch in range( from_idx , to_idx + 1 , 1 ):
            path = local_path + "/dev1num3model_for_epoch{Epoch}.pth".format( Epoch = epoch )
            tmp_dict = torch.load(path)
            state_dict_list.append(tmp_dict)
            # print( path )
        else :
            cudanet = Classifier( kernel_sz = 9 )
            cudanet.load_state_dict(tmp_dict)
            cudanet.to(device)
            cudanet.eval()

```